

AI Planning

Artificial Intelligence

Reading

- Required reading
 - Sections 11.1 — 11.4
- Recommended reading
 - AIMA Section 10.3: Actions, Situations, and Events
 - Chapter 11 entirely

Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- Algorithms

Background

- Focus
 - The focus here is deterministic planning
 - Environment is fully observable
 - Results of actions is deterministic
 - Relaxing the above requires dealing with uncertainty
 - Problem types: sensorless, contingency, exploration
- Planning ‘communities’ in AI
 - Logic-based: Reasoning About Actions & Change
 - Less formal representations: Classical AI Planning
 - Uncertainty (UAI): Graphical Models such as
 - Markov Decision Processes (MDP), Partially Observable MDPs, etc.
- AI Planning is **not** MRP (Material Requirements Planning)


Actions, events, and change

- Planning requires a representation of time
 - to express & reason about sequences of actions
 - to express the effects of actions on the world
- Propositional Logic
 - does not offer a representation for time
 - Each action description needs to be repeated for each step
- Situation Calculus (AIMA Section 10.3)
 - Is based on FOL
 - Each time step is a ‘situation’
 - Allows to represent plans and reason about actions & change

Situation Calculus: Ontology


- Situations
- Fluents
- Atemporal (or eternal) predicates & functions

AIMA Section 10.3



Situation Calculus: Ontology

- Situations
 - Initial state: S_0
 - A function $Result(a.s)$ gives the situation resulting from applying action a in situation s
- Fluents
 - Functions & predicates whose truth values can change from one situation to the other
 - Example: $\neg Holding(G_1, S_0)$
- Atemporal (or eternal) predicates and functions
 - Example: $Gold(G_1)$, $LeftLegOf(Wumpus)$



Situation Calculus

- Sequence of actions
 - $\text{Result}([], s) = s$
 - $\text{Result}([a|\text{seq}], s) = \text{Result}(\text{seq}, \text{Result}(a, s))$
- Projection task
 - Deducing the outcome of a sequence of actions
- Planning task
 - Find a sequence of actions that achieves a desired effect

Example: Wumpus World

- Fluents
 - $\text{At}(o,p,s), \text{Holding}(o,s)$
- Agent is in [1,1], gold is in [1,2]
 - $\text{At}(\text{Agent},[1,1],S_0) \wedge \text{At}(G_1,[1,2],S_0)$
- In S_0 , we also need to have:
 - $\text{At}(o,x,S_0) \Leftrightarrow [(o=\text{Agent}) \wedge x=[1,1]] \vee [(o=G_1) \wedge x=[1,2]]$
 - $\neg \text{Holding}(o,S_0)$
 - $\text{Gold}(G_1) \wedge \text{Adjacent}([1,1],[1,2]) \wedge \text{Adjacent}([1,2],[1,1])$
- The query is:
 - $\exists \text{seq } \text{At}(G_1,[1,1],\text{Result}(\text{seq},S_0))$
- The answer is
 - $\text{At}(G_1,[1,1],\text{Result}(\text{Go}([1,1],[1,2]),\text{Grab}(G_1),\text{Go}([1,2],[1,1]),S_0))$

Importance of Situation Calculus

- Historical note
 - Situation Calculus was the first attempt to formalizing planning in FOL
 - Other formalisms include Event Calculus
 - The area of using logic for planning is informally called in the literature “Reasoning About Action & Change”
- Highlighted three important problems
 1. Frame problem
 2. Qualification problem
 3. Ramification problem

‘Famous’ Problems

- Frame problem
 - Representing all things that stay the same from one situation to the next
 - Inferential and representational
- Qualification problem
 - Defining the circumstances under which an action is guaranteed to work
 - Example: what if the gold is slippery or nailed down, etc.
- Ramification problem
 - Proliferation of implicit consequences of actions as actions may have secondary consequences
 - Examples: How about the dust on the gold?

Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- **Representation language**
- Algorithms

Planning Languages

- Languages must represent..
 - States
 - Goals
 - Actions
- Languages must be
 - Expressive for ease of representation
 - Flexible for manipulation by algorithms

General language features

- Representation of states
 - Decompose the world in logical conditions and represent a state as a conjunction of positive literals.
 - Propositional literals: $Poor \wedge Unknown$
 - FO-literals (ground and function-free): $At(Plane1, Melbourne) \wedge At(Plane2, Sydney)$
 - Closed world assumption
- Representation of goals
 - Partially specified state and represented as a conjunction of positive ground literals
 - $Poor \wedge Unknown \wedge At(P2, Tahiti)$
 - A goal is satisfied if the state contains all literals in goal.

General language features

- Representations of actions

- Action = PRECOND + EFFECT

Action(Fly(p,from, to),

PRECOND: $At(p,from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT: $\neg At(p,from) \wedge At(p,to)$

= action schema (p, from, to: need to be instantiated)

- Action name and parameter list
- Precondition (conj. of function-free positive literals)
- Effect (conj of function-free literals and P is True and not P is false)
- Add-list vs delete-list in Effect

State Representation

- A state is represented with a conjunction of positive literals
- Using
 - Logical Propositions: $Poor \wedge Unknown$
 - FOL literals: $At(Plane1, OMA) \wedge At(Plan2, JFK)$
- FOL literals must be ground & function-free
 - **Not allowed**: $At(x, y)$ or $At(Father(Fred), Sydney)$
- Closed World Assumption
 - What is not stated are assumed false

Goal Representation

- Goal is a partially specified state
- A proposition satisfies a goal if it contains all the atoms of the goal and possibly others..
 - Example: $\text{Rich} \wedge \text{Famous} \wedge \text{Miserable}$ satisfies the goal $\text{Rich} \wedge \text{Famous}$

Action Representation

- Action Schema

- Action name
- Preconditions
- Effects

$At(WHI, LNK), Plane(WHI),$
 $Airport(LNK), Airport(OHA)$

$Fly(WHI, LNK, OHA)$

$At(WHI, OHA), \neg At(WHI, LNK)$

- Example

Action($Fly(p, from, to)$,

PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT: $\neg At(p, from) \wedge At(p, to)$)

- Sometimes, Effects are split into ADD list and DELETE list

Applying an Action

- Find a substitution list θ for the variables
 - of all the precondition literals
 - with (a subset of) the literals in the current state description
- Apply the substitution to the propositions in the effect list
- Add the result to the current state description to generate the new state
- Example:
 - Current state: $\text{At}(\text{P1}, \text{JFK}) \wedge \text{At}(\text{P2}, \text{SFO}) \wedge \text{Plane}(\text{P1}) \wedge \text{Plane}(\text{P2}) \wedge \text{Airport}(\text{JFK}) \wedge \text{Airport}(\text{SFO})$
 - It satisfies the precondition with $\theta = \{p/\text{P1}, \text{from}/\text{JFK}, \text{to}/\text{SFO}\}$
 - Thus the action $\text{Fly}(\text{P1}, \text{JFK}, \text{SFO})$ is applicable
 - The new current state is: $\text{At}(\text{P1}, \text{SFO}) \wedge \text{At}(\text{P2}, \text{SFO}) \wedge \text{Plane}(\text{P1}) \wedge \text{Plane}(\text{P2}) \wedge \text{Airport}(\text{JFK}) \wedge \text{Airport}(\text{SFO})$


Languages for Planning Problems

- STRIPS
 - Stanford Research Institute Problem Solver
 - Historically important
- ADL
 - Action Description Languages
 - See Table 11.1 for STRIPS versus ADL
- PDDL
 - Planning Domain Definition Language
 - Revised & enhanced for the needs of the International Planning Competition
 - Currently [version 3.1](#)

Expressiveness and extensions

- STRIPS is simplified
 - Important limit: function-free literals
 - Allows for propositional representation
 - Function symbols lead to infinitely many states and actions
- Recent extension: Action Description language (ADL)
 - Action(Fly(*p*:Plane, from: Airport, to: Airport),
 - PRECOND: $At(p, from) \wedge (from \neq to)$
 - EFFECT: $\neg At(p, from) \wedge At(p, to)$
- Standardization : Planning domain definition language (PDDL)

STRIPS



STRIPS - Box in the Room Example

//Initial state:

INROOM(ROBOT, R1) \wedge CONNECTS(D1, R1, R2) \wedge CONNECTS(D2, R2, R3) \wedge BOX(BOX1),
INROOM(BOX1, R2)

//Goal:

BOX(BOX1) \wedge INROOM(BOX1, R1)

//Actions:

GOTHRU(d, r1, r2) //Robot goes from room r1 to room r2, through door d

Preconditions: INROOM(ROBOT, r1) \wedge CONNECTS(d, r1, r2)

Postconditions: NOT INROOM(ROBOT, r1) \wedge INROOM(ROBOT, r2)

PUSHTHRU(b, d, r1, r2) //Robot pushes box b through door d from room r1 to room r2

Preconditions: INROOM(b, r1) \wedge INROOM(ROBOT, r1) \wedge CONNECTS(d, r1, r2)

Postconditions: NOT INROOM(ROBOT, r1) \wedge NOT INROOM(b, r1) \wedge INROOM(ROBOT, r2) \wedge INROOM(b, r2)

Here is the solution (sequence of actions to achieve the goal, while starting from initial state) for described example:

1. GOTHRU(D1,R1,R2)
2. PUSHTHRU(BOX1,D1,R2,R1).

Example: Air Cargo

- See Figure 11.2
- Initial state
- Goal State
- Actions: Load, Unload, Fly

Example: air cargo transport

$Init(At(C1, SFO) \wedge At(C2, JFK) \wedge At(P1, SFO) \wedge At(P2, JFK) \wedge Cargo(C1) \wedge$
 $Cargo(C2) \wedge Plane(P1) \wedge Plane(P2) \wedge Airport(JFK) \wedge Airport(SFO))$
 $Goal(At(C1, JFK) \wedge At(C2, SFO))$

$Action(Load(c, p, a))$

PRECOND: $At(c, a) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $\neg At(c, a) \wedge In(c, p)$

$Action(Unload(c, p, a))$

PRECOND: $In(c, p) \wedge At(p, a) \wedge Cargo(c) \wedge Plane(p) \wedge Airport(a)$

EFFECT: $At(c, a) \wedge \neg In(c, p)$

$Action(Fly(p, from, to))$

PRECOND: $At(p, from) \wedge Plane(p) \wedge Airport(from) \wedge Airport(to)$

EFFECT: $\neg At(p, from) \wedge At(p, to)$

[$Load(C1, P1, SFO)$, $Fly(P1, SFO, JFK)$, $Load(C2, P2, JFK)$, $Fly(P2, JFK, SFO)$, ...]

Example: Spare Tire Problem

- See Figure 11.3
- Initial State
- Goal State
- Actions:
 - *Remove(Spare, Trunk), Remove(Flat, Axle)*
 - *PutOn(Spare, Axle)*
 - *LeaveOvernight*
- Note
 - the negated precondition $\neg At(Flat, Axle)$ not allowed in STRIPS.
 - Could be easily replaced with *Clear(Axle)*, adding one more predicate to the language

Example: Spare tire problem

Init(*At*(Flat, Axle) \wedge *At*(Spare, trunk))

Goal(*At*(Spare, Axle))

Action(*Remove*(Spare, Trunk)

PRECOND: *At*(Spare, Trunk)

EFFECT: \neg *At*(Spare, Trunk) \wedge *At*(Spare, Ground))

Action(*Remove*(Flat, Axle)

PRECOND: *At*(Flat, Axle)

EFFECT: \neg *At*(Flat, Axle) \wedge *At*(Flat, Ground))

Action(*PutOn*(Spare, Axle)

PRECOND: *At*(Spare, Ground) \wedge \neg *At*(Flat, Axle)

EFFECT: *At*(Spare, Axle) \wedge \neg *At*(Spare, Ground))

Action(*LeaveOvernight*

PRECOND:

EFFECT: \neg *At*(Spare, Ground) \wedge \neg *At*(Spare, Axle) \wedge \neg *At*(Spare, trunk) \wedge \neg *At*(Flat, Ground) \wedge \neg *At*(Flat, Axle))



STRIPS?

Example: Blocks World

- See Fig 11.4
- Initial state
- Goal
- Actions:
 - *Move(b,x,y)*
 - *MoveToTable(b,x)*

Example: Blocks world

$Init(On(A, Table) \wedge On(B, Table) \wedge On(C, Table) \wedge Block(A) \wedge Block(B) \wedge Block(C) \wedge Clear(A) \wedge Clear(B) \wedge Clear(C))$

$Goal(On(A, B) \wedge On(B, C))$

$Action(Move(b, x, y))$

PRECOND: $On(b, x) \wedge Clear(b) \wedge Clear(y)$

$\wedge Block(b) \wedge (b \neq x) \wedge (b \neq y) \wedge (x \neq y)$


EFFECT: $On(b, y) \wedge Clear(x) \wedge \neg On(b, x) \wedge \neg Clear(y)$

$Action(MoveToTable(b, x))$

PRECOND: $On(b, x) \wedge Clear(b) \wedge Block(b) \wedge (b \neq x)$

EFFECT: $On(b, Table) \wedge Clear(x) \wedge \neg On(b, x)$

Spurious actions: $Move(B, C, C)$



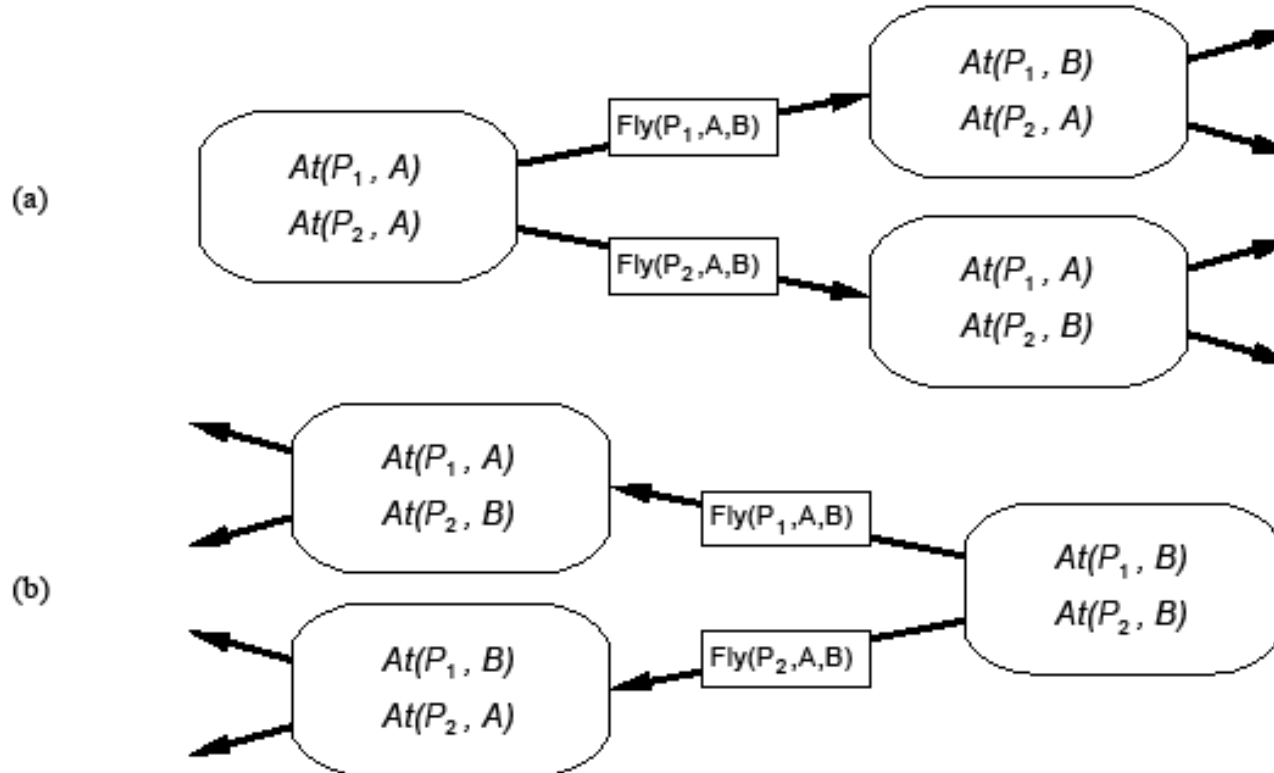
Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- **Planning Algorithms**
 - **State-Space Search**
 - Partial-Order Planning (POP)
 - Planning Graphs (GRAPHPLAN)
 - SAT Planners

State-Space Search (1)

- Search the space of states (first chapters)
 - Initial state, goal test, step cost, etc.
 - Actions are the transitions between state
- Actions are invertible (why?)
 - Move forward from the initial state: Forward State-Space Search or Progression Planning
 - Move backward from goal state: Backward State-Space Search or Regression Planning

State-Space Search (2)

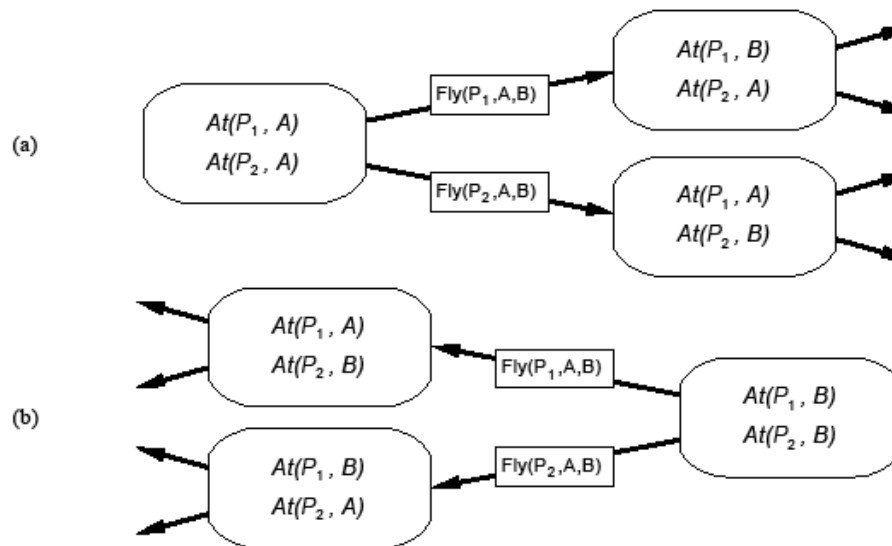


State-Space Search (3)

- Remember that the language has no functions symbols
- Thus number of states is finite
- And we can use any complete search algorithm (e.g., A^*)
 - We need an admissible heuristic
 - The solution is a path, a sequence of actions: **total-order planning**
- Problem: Space and time complexity
 - STRIPS-style planning is PSPACE-complete unless actions have
 - only positive preconditions and
 - only one literal effect

SRIPS in State-Space Search

- STRIPS representation makes it easy to focus on ‘relevant’ propositions and
 - Work backward from goal (using EFFECTS)
 - Work forward from initial state (using PRECONDITIONS)
 - Facilitating bidirectional search



Relevant Action

- An action is relevant
 - In Progression planning when its preconditions match a subset of the current state
 - In Regression planning, when its effects match a subset of the current goal state

Consistent Action

- The purpose of applying an action is to 'achieves a desired literal'
- We should be careful that the action does not undo a desired literal (as a side effect)
- A consistent action is an action that does not undo a desired literal

Backward State-Space Search

- Given
 - A goal G description
 - An action A that is relevant and consistent
- Generate a predecessor state where
 - Positive effects (literals) of A in G are deleted
 - Precondition literals of A are added unless they already appear
 - Substituting any variables in A 's effects to match literals in G
 - Substituting any variables in A 's preconditions to match substitutions in A 's effects
- Repeat until predecessor description matches initial state

Heuristic to Speed up Search

- We can use A^* , but we need an admissible heuristic
 1. Divide-and-conquer: sub-goal independence assumption
 - Problem relaxation by removing
 2. ... all preconditions
 3. ... all preconditions and negative effects
 4. ... negative effects only: Empty-Delete-List

1. Subgoal Independence Assumption

- The cost of solving a conjunction of subgoals is the sum of the costs of solving each subgoal independently
- Optimistic
 - Where subplans interact negatively
 - Example: one action in a subplan delete goal achieved by an action in another subplan
- Pessimistic (not admissible)
 - Redundant actions in subplans can be replaced by a single action in merged plan

2. Problem Relaxation: Removing Preconditions

- Remove preconditions from action descriptions
 - All actions are applicable
 - Every literal in the goal is achievable in one step
- Number of steps to achieve the conjunction of literals in the goal is equal to the number of unsatisfied literals
- Alert
 - Some actions may achieve several literals
 - Some action may remove the effect of another action

3. Remove Preconditions & Negative Effects

- Considers only positive interactions among actions to achieve multiple subgoals
- The minimum number of actions required is the sum of the union of the actions' positive effects that satisfy the goal
- The problem is reduced to a set cover problem, which is NP-hard
 - Approximation by a greedy algorithm cannot guarantee an admissible heuristic

4. Removing Negative Effects (Only)

- Remove all negative effects of actions (no action may destroy the effects of another)
- Known as the Empty-Delete-List heuristic
- Requires running a simple planning algorithm
- Quick & effective
- Usable in progression or regression planning

Outline

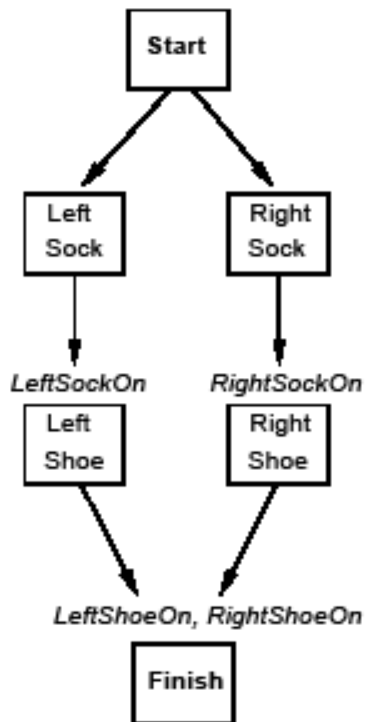
- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- **Planning Algorithms**
 - State-Space Search
 - **Partial-Order Planning (POP)**
 - Planning Graphs (GRAPHPLAN)
 - SAT Planners

Partial Order Planning (POP)

- State-space search
 - Yields totally ordered plans (linear plans)
- POP
 - Works on subproblems independently, then combines subplans
 - Example
 - Goal(RightShoeOn \wedge LeftShoeOn)
 - Init()
 - *Action*(RightShoe, PRECOND: RightSockOn, EFFECT: RightShoeOn)
 - *Action*(RightSock, EFFECT: RightSockOn)
 - *Action*(LeftShoe, PRECOND: LeftSockOn, EFFECT: LeftShoeOn)
 - *Action*(LeftSock, EFFECT: LeftSockOn)

POP Example & its linearization

Partial Order Plan:



Total Order Plans:



Components of a Plan

1. A set of **actions**
2. A set of **ordering constraints**
 - A $[?][?]$ B reads “A before B” but not necessarily immediately before B
 - Alert: caution to cycles A $[?][?]$ B and B $[?][?]$ A
3. A set of **causal links** (protection intervals) between actions
 - A \xrightarrow{p} B reads “A achieves p for B” and p must remain true from the time A is applied to the time B is applied
 - Example “RightSock $\xrightarrow{\text{RightSockOn}}$ RightShoe”
4. A set of **open preconditions**
 - Planners work to reduce the set of open preconditions to the empty set w/o introducing contradictions

Consistent Plan (POP)

- Consistent plan is a plan that has
 - No cycle in the ordering constraints
 - No conflicts with the causal links
- Solution
 - Is a consistent plan with no open preconditions
- To solve a conflict between a causal link $A \xrightarrow{p} B$ and an action C (that **clobbers**, threatens the causal link), we force C to occur outside the “protection interval” by adding
 - the constraint $C \boxed{?} \boxed{?} A$ (**demoting** C) or
 - the constraint $B \boxed{?} \boxed{?} C$ (**promoting** C)

Setting up the PoP

- Add dummy states

- Start

- Has no preconditions
 - Its effects are the literals of the initial state

- Finish

- Its preconditions are the literals of the goal state
 - Has no effects

Start
Literal_a, Literal_b, ...

Literal₁, Literal₂, ...
Finish

- Initial Plan:

- Actions: {Start, Finish}

- Ordering constraints: {Start [?] [?] Finish}

- Causal links: {}

- Open Preconditions: {LeftShoeOn, RightShoeOn}

Start

LeftShoeOn, RightShoeOn

Finish

POP as a Search Problem

- The successor function arbitrarily picks one open precondition p on an action B
- For every possible consistent action A that achieves p
 - It generates a successor plan adding the causal link $A \xrightarrow{p} B$ and the ordering constraint $A \square \square B$
 - If A was not in the plan, it adds $\text{Start} \square \square A$ and $A \square \square \text{Finish}$
 - It resolves all conflicts between
 - the new causal link and all existing actions
 - between A and all existing causal links
 - Then it adds the successor states for combination of resolved conflicts
- It repeats until no open precondition exists

Example of POP: Flat tire problem

- See problem description in Fig 11.7 page 391

Start
At(Spare,Trunk), At(Flat,Axle)

- Only one open precondition
- Only 1 applicable action

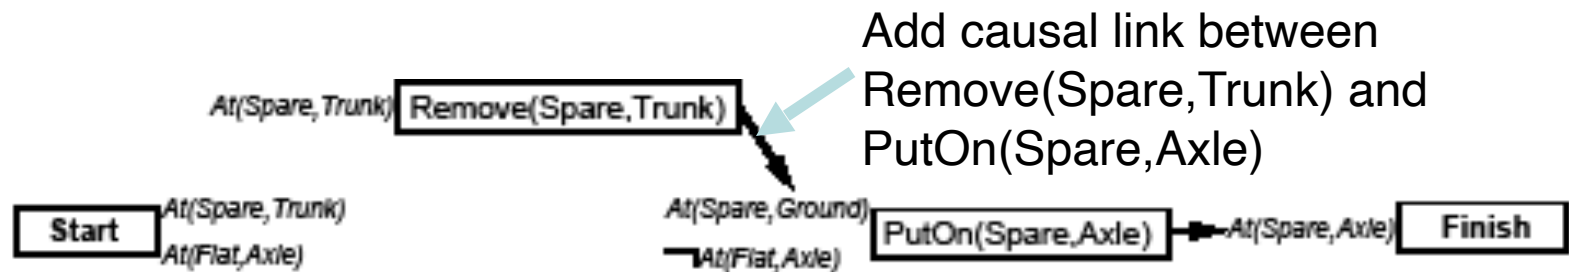
At(Spare,Ground), \neg At(Flat,Axle)

PutOn(Spare,Axle)

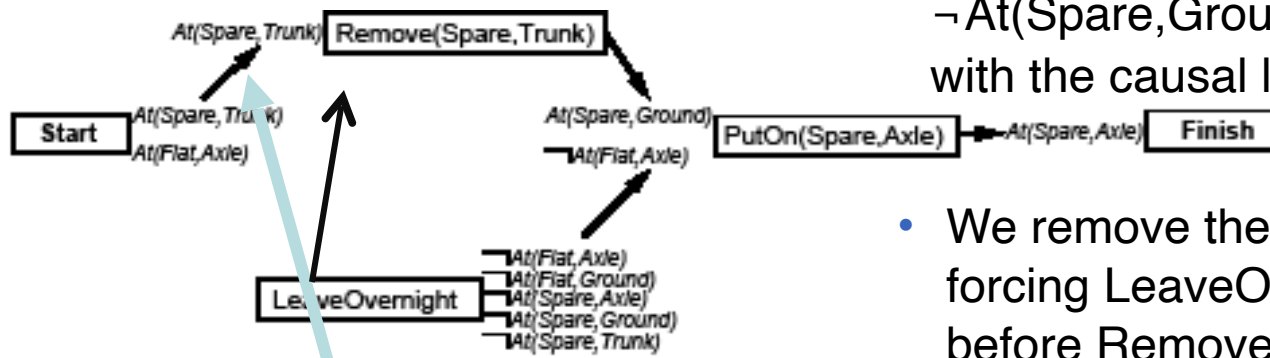
- Pick up At(Spare,Ground)
- Choose only applicable action
Remove(Spare,Trunk)

At(Spare,Axle)


Finish



- Pick up $\neg At(Flat, Axle)$
- There are 2 applicable actions: $LeaveOvernight$ and $Remove(Flat, Axle)$
- Choose $LeaveOvernight$



- Conflicts with effects of $Remove(Spare, Trunk)$
- The only way to resolve the conflict is to undo $LeaveOvernight$ use the action $Remove(Flat, Axle)$



- This time, we choose Remove(Flat,Axle)
- Pick up At(Spare,Trunk) and choose Start to achieve it
- Pick up At(Flat,Axle) and choose Start to achieve it.
- We now have a complete consistent partially ordered plan

POP Algorithm (1)

- Backtrack when fails to resolve a threat or find an operator
- Causal links
 - Recognize when to abandon a doomed plan without wasting time expanding irrelevant part of the plan
 - allow early pruning of inconsistent combination of actions
- When actions include variables, we need to find appropriate substitutions
 - Typically we try to delay commitments to instantiating a variable until we have no other choice (least commitment)
- POP is sound, complete, and systematic (no repetition)

POP Algorithm (2)

- Decomposes the problem (advantage)
- But does not represent states explicitly: it is hard to design heuristic to estimate distance from goal
 - Example: Number of open preconditions – those that match the effects of the start node. Not perfect (same problems as before)
- A heuristic can be used to choose which plan to refine (which precondition to pick-up):
 - Choose the most-constrained precondition, the one satisfied by the least number of actions. Like in CSPs!
 - When no action satisfies a precondition, backtrack!
 - When only one action satisfies a precondition, pick up the precondition.

Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- **Planning Algorithms**
 - State-Space Search
 - Partial-Order Planning (POP)
 - **Planning Graphs (GRAPHPLAN)**
 - SAT Planners

Planning Graph

- Is special data structure used for
 1. Deriving better heuristic estimates
 2. Extract a solution to the planning problem: GRAPHPLAN algorithm
- Is a sequence $\langle S_0, A_0, S_1, A_1, \dots, S_i \rangle$ of levels
 - Alternating state levels & action levels
 - Levels correspond to time stamps
 - Starting at initial state
 - State level is a set of (propositional) literals
 - All those literals that could be true at that level
 - Action level is a set of (propositionalized) actions
 - All those actions whose preconditions appear in the state level (ignoring all negative interactions, etc.)
- Propositionalization may yield combinatorial explosion in the presence of a large number of objects

Focus

- Building the Planning Graph
- Using it for Heuristic Estimation
- Using it for generating the plan

Example of a Planning Graph (1)

Init(Have(Cake))

Goal(Have(Cake) \wedge Eaten(Cake))

Action(Eat(Cake))

Precond: Have(Cake)

Effect: \neg Have(Cake) \wedge Eaten(Cake))

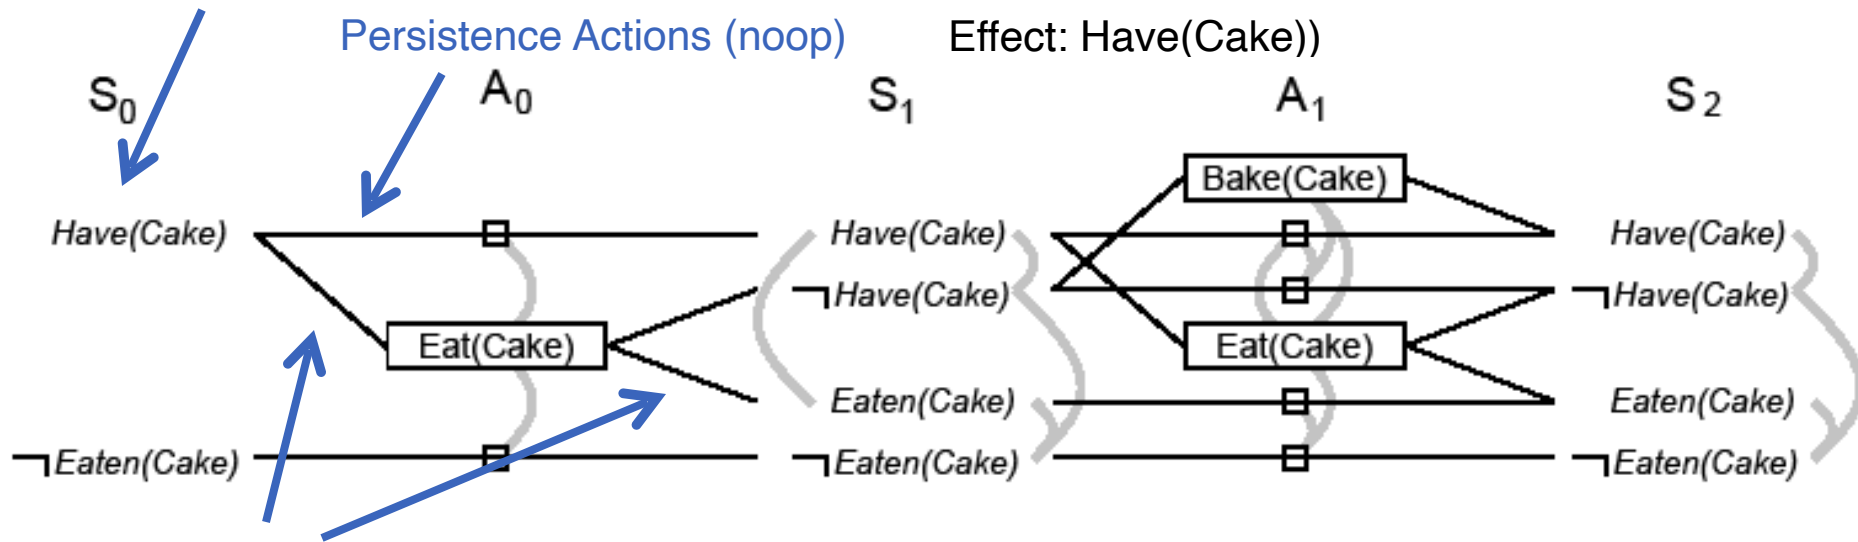
Action(Bake(Cake))

Precond: \neg Have(Cake)

Effect: Have(Cake))

Propositions true
at the initial state


Persistence Actions (noop)



Action is connected to its
preconds & effects


Example of a Planning Graph (2)

- At each state level, list all literals that may hold at that level
- At each action level, list all noops & all actions whose preconditions may hold at previous levels
- Repeat until plan 'levels off,' no new literals appears ($S_i = S_{i+1}$)
- Building the Planning Graph is a polynomial process
- Add (binary) mutual exclusion (mutex) links between conflicting actions and between conflicting literals




Mutex Links between Actions

1. **Inconsistent effects:** one action negates an effect of another
 - Eat(Cake) & noop of Have(Cake) disagree on effect Have(Cake)
2. **Interference:** An action effect negates the precondition of another
 - Eat(Cake) negates precondition of the noop of Have(Cake):
3. **Competing needs:** A precondition on an action is mutex with the precondition of another
 - Bake(Cake) & Eat(Cake): compete on Have(Cake) precondition



Mutex Links between Literals

1. Two literals are negation of each other
2. **Inconsistent support:** Each pair of actions that can achieve the two literals is mutex. Examples:
 - In S1, Have(Cake) & Eaten(Cake) are mutex
 - In S2, they are not because Bake(Cake) & the noop of Eaten(Cake) are not mutex



Focus

- Building the Planning Graph
- **Using it for Heuristic Estimation**
 - Planning graph as a relaxation of original problem
 - Easy to build (compute)
- Using it for generating the plan

Planning Graph for Heuristic Estimation

- A literal that does not appear in the final level cannot be achieved by any plan
 - State-space search: Any state containing an unachievable literal has cost $h(n)=\infty$
 - POP: Any plan with an unachievable open condition has cost $h(n)=\infty$
- The estimate cost of any goal literal is the first level at which it appears
 - Estimate is admissible for individual literals
 - Estimate can be improved by serializing the graph (serial planning graph: one action per level) by adding mutex between all actions in a given level
- The estimate of a conjunction of goal literals
 - Three heuristics: max level, level sum, set level

Estimate of Conjunction of Goal Literals

- Max-level
 - The largest level of a literal in the conjunction
 - Admissible, not very accurate
- Level sum
 - Under subgoal independence assumption, sums the level costs of the literals
 - Inadmissible, works well for largely decomposable problems
- Set level
 - Finds the level at which all literals appear w/o any pair of them being mutex
 - Dominates max-level, works extremely well on problems where there is a great deal of interaction among subplans

Focus

- Building the Planning Graph
- Using it for Heuristic Estimation
- **Using it for generating the plan**
 - GraphPlan algorithm [Blum & Furst, 95]

GRAPHPLAN algorithm

GRAPHPLAN(*problem*) **returns** *solution* or *failure*

graph \leftarrow INITIALPLANNINGGRAPH(*problem*)

goals \leftarrow GOALS[*problem*]

loop do

if *goals* all non-mutex in last level of graph **then do**

solution \leftarrow EXTRACTSOLUTION(*graph*,*goals*,LENGTH(*graph*))

if *solution* \neq *failure* **then return** *solution*

else if NOSOLUTIONPOSSIBLE(*graph*) **then return** *failure*

graph \leftarrow EXPANDGRAPH (*graph*,*problem*)

- Two main stages
 1. Extract solution
 2. Expand the graph

Example of GRAPHPLAN Execution (1)

- $At(Spare, Axle)$ is not in S_0
- No need to extract solution
- Expand the plan

S_0
 $At(Spare, Trunk)$

$At(Flat, Axle)$


$\neg At(Spare, Axle)$

$\neg At(Flat, Ground)$

$\neg At(Spare, Ground)$

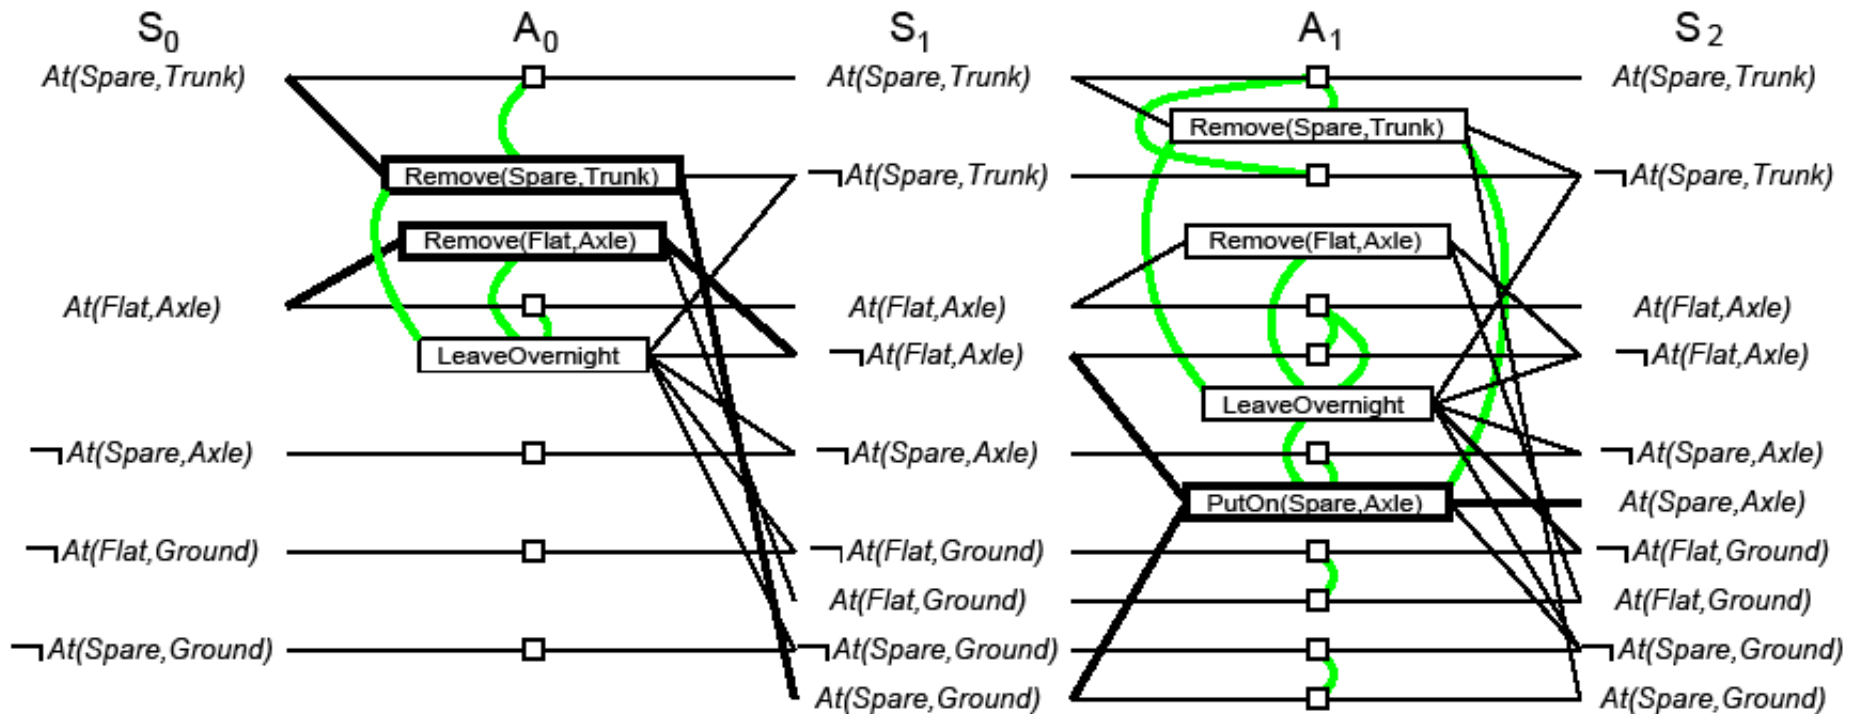
Example of GRAPHPLAN Execution (2)

- Three actions are applicable
- 3 actions and 5 noops are added
- Mutex links are added
- $At(Spare, Axle)$ still not in S_1
- Plan is expanded




Example of GRAPHPLAN Execution (3)

- Illustrates well mutex links: inconsistent effects, interference, competing needs, inconsistent support



Solution Extraction (Backward)

1. Solve a Boolean CSP: Variables are actions, domains are {0=out of plan, 1=in plan}, constraints are mutex
2. Search problem from last level backward



Backtrack Search for Solution Extraction

- Starting at the highest fact level
 - Each goal is put in a goal list for the current fact layer
 - Search iterates thru each fact in the goal list trying to find an action to support it which is not mutex with any other chosen action
 - When an action is chosen, its preconditions are added to the goal list of the lower level
 - When all facts in the goal list of the current level have a consistent assignment of actions, the search moves to the next level
- Search backtracks to the previous level when it fails to assign an action to each fact in the goal list at a given level
- Search succeeds when the first level is reached.

Termination of GRAPHPLAN

- GRAPHPLAN is guaranteed to terminate
 - Literal increase monotonically
 - Actions increase monotonically
 - Mutexes decrease monotonically
- A solution is guaranteed not to exist when
 - The graph levels off with all goals present & non-mutex, and
 - **EXTRACTSOLUTION** fails to find solution

Optimality of GRAPHPLAN

- The plans generated by GRAPHPLAN
 - Are optimal in the number of steps needed to execute the plan
 - Not necessarily optimal in the number of actions in the plan (GRAPHPLAN produces partially ordered plans)

Outline

- Background
 - Situation Calculus
 - Frame, qualification, & ramification problems
- Representation language
- Planning Algorithms
 - State-Space Search
 - Partial-Order Planning (POP)
 - Planning Graphs (GRAPHPLAN)
 - SAT Planners